

Sycophancy in Hierarchical versus Flat Multi-Agent Systems*

Gautham Velupally

University of Illinois
Urbana-Champaign
Urbana, IL, USA
gkv4@illinois.edu

Pranav Nagarajan

University of Illinois
Urbana-Champaign
Urbana, IL, USA
pranavn6@illinois.edu

Shanmukh Upadhyayula

University of Illinois
Urbana-Champaign
Urbana, IL, USA
supad83@illinois.edu

Sanjit Kalle

University of Illinois
Urbana-Champaign
Urbana, IL, USA
skal1e4@illinois.edu

Abstract

Multi-agent systems (MAS) built on large language models are increasingly deployed for complex reasoning tasks in finance, policy, and medicine. We investigate whether hierarchical MAS topologies systematically amplify hallucination propagation compared to flat topologies, a phenomenon we term the *Yes-Man collapse*. In a controlled experiment using 21 Gemini 2.5 Flash agents performing a geopolitical market-prediction task, we inject an identical hallucination into one agent in each topology: the top-level orchestrator in the hierarchy, and a random peer in the flat group. In the hierarchical condition, 100% of agents adopt the hallucinated stance by Turn 3, yielding a sycophancy effect size of $\Delta^2 = 0.80$. In the flat condition, the hallucination is contained to the single injected peer ($\Delta^2 = 0.048$). All 84 failed agent-turns in the hierarchy are classified as *planning errors*, agents abandoned correct evidence-based reasoning to defer to authority, despite explicit anti-sycophancy instructions in every agent’s persona. These results demonstrate that topology is a safety-critical design choice for MAS deployment and that constitutional instructions alone are insufficient to prevent authority-driven cascade failures.

1 Introduction

Large language models (LLMs) exhibit sycophantic behavior. They tend to agree with users’ stated views even when those views are incorrect (Perez et al., 2023; Sharma et al., 2024). This phenomenon has been extensively studied in single-agent, single-turn settings, where it manifests as models parroting user opinions, conceding to incorrect premises, and generating responses that prioritize agreeableness over accuracy (Ranaldi and Freitas, 2024).

As LLMs are increasingly composed into multi-agent systems (MAS), teams of AI agents that collaborate through structured communication to solve complex tasks (Qian et al., 2024; Park et al., 2023; Du et al., 2023), a natural question emerges. *Does sycophancy propagate between agents, and does the communication topology affect the rate and extent of propagation?*

This question has immediate practical consequences. MAS are being deployed for high-stakes reasoning tasks including financial analysis, policy simulation, and medical decision support. Many production systems adopt hierarchical topologies that mirror human organizational structures, where a lead agent synthesizes reports from subordinate agents, who in turn manage their own sub-teams (Qian et al., 2024; Hong et al., 2024). If hierarchical structure amplifies sycophancy, causing a single hallucination at the top to cascade into unanimous system failure, then topology becomes a safety-critical design choice, not merely an engineering convenience.

The analogy to human organizations is instructive. Decades of organizational behavior research have documented how authority gradients suppress dissent in corporate hierarchies (Janis, 1972; Asch, 1951). The Challenger disaster, the 2008 financial crisis, and numerous corporate frauds share a common structural feature, where correct information existed at lower levels of the hierarchy but was filtered, suppressed, or overridden by authority signals from above. We ask whether LLM-based multi-agent systems, despite lacking human social motivations like career concerns or fear of retaliation, reproduce these same dynamics through the structure of their communication topology alone.

We present the first controlled measurement of authority-gradient sycophancy propagation across MAS topologies, guided by two research questions.

*Code at <https://github.com/VelupallyG/mas-sycophancy>.

- **RQ1.** Does a hallucination injected at the top of a hierarchy fully propagate down, and how fast?
- **RQ2.** Can a lower-ranked correct agent shift consensus against a wrong authority?

Our contributions are as follows.

1. A controlled experimental framework that isolates topology as the independent variable by holding constant the model, evidence, agent count, and hallucination content across flat and hierarchical conditions.
2. Novel metrics, Δ^2 (sycophancy effect size), Turn of Flip (ToF), and Number of Flips (NoF), that extend single-agent sycophancy probes (Perez et al., 2023; Sharma et al., 2024) to inter-agent cascade dynamics.
3. Empirical evidence that a single hallucination injected at the top of a 21-agent hierarchy causes 100% system capture within 3 turns, while the identical hallucination in a flat topology is contained to the injector (4.8% failure rate).
4. A linguistic analysis showing that agents transition from explicit authority deference to implicit adoption of the hallucinated stance, manufacturing post-hoc justifications for incorrect predictions.

2 Background and Related Work

2.1 Multi-Agent Systems and Simulation

Park et al. (2023) demonstrated that LLM-powered generative agents simulate believable human behavior in open-ended environments, exhibiting emergent information diffusion, relationship formation, and group coordination without explicit social rules.

Vezhnevets et al. (2023) introduced the Concordia framework for generative agent-based modeling, which grounds agent actions in physical, social, or digital spaces. Concordia provides the component-based architecture we build upon, where agents are constructed by composing modular components (memory, observation, instructions, action) rather than through monolithic prompts. This modularity enables controlled experimentation, and we vary topology and rank components while holding all other agent internals constant. Prior multi-agent debate work

(Du et al., 2023), alongside swarm-synthesis approaches such as MiroFish (MiroFish Research, 2024), establishes that deliberation across agents can outperform single-agent analysis on complex reasoning and forecasting, making the topology question practically relevant rather than academic.

2.2 Hierarchical and Role-Based MAS Architectures

Several influential systems have adopted hierarchical topologies for LLM-based multi-agent collaboration. Qian et al. (2024) introduced ChatDev, where specialized agents (CEO, CTO, programmer, tester) collaborate through structured dialogue chains along a software development waterfall. Agents are assigned distinct roles with explicit authority relationships, where the CEO issues requirements, the CTO makes architectural decisions, and programmers implement. This mirrors real corporate structures and established that hierarchical communication graphs can coordinate complex multi-step tasks.

Hong et al. (2024) proposed MetaGPT, which encodes human-like Standard Operating Procedures (SOPs) into multi-agent collaboration, with agents assigned to specific roles in a workflow pipeline.

These systems demonstrate the practical utility of hierarchical MAS but do not study the failure modes introduced by authority gradients. Our work fills this gap by empirically testing what happens when the agent at the top of such a hierarchy confidently asserts an incorrect conclusion.

2.3 Sycophancy in Single LLMs

Perez et al. (2023) introduced auto-generated evaluations demonstrating that RLHF-trained models exhibit sycophantic behavior that worsens with scale, larger models are more likely to agree with user-stated opinions regardless of factual accuracy. Sharma et al. (2024) extended this finding across five frontier RLHF assistants, showing that models concede to incorrect user assertions and tracing the root cause to preference data that rewards convincing-sounding responses over correct ones. Their analysis revealed that sycophancy is not a superficial behavior but is deeply embedded in the reward signal, since human evaluators systematically prefer responses that agree with them even when those responses are wrong.

Ranaldi and Freitas (2024) further characterized the conditions under which LLMs contradict hu-

man users versus capitulate, finding that contradiction is more likely when the model’s confidence is very high and the user assertion is clearly factual rather than opinion-based. [Wei et al. \(2024\)](#) showed that simple synthetic training data can reduce sycophancy, suggesting the behavior is learned and potentially correctable at the training level.

Critically, all of this work operates in single-agent, single-turn settings, with one model interacting with one user in one exchange. Our work extends sycophancy measurement to inter-agent propagation across multiple deliberation turns, where the “user” exerting pressure is itself another LLM agent with structural authority.

2.4 Multi-Agent Debate and Deliberation

A growing body of work explores whether multi-agent interaction improves LLM reasoning. [Du et al. \(2023\)](#) demonstrated that multi-agent debate, where multiple LLM instances critique and refine each other’s responses, improves factuality and reasoning over single-agent baselines on arithmetic, strategic reasoning, and factual QA tasks.

However, the benefits of multi-agent deliberation are not universal. [Wang et al. \(2024\)](#) questioned whether multi-agent discussion reliably improves reasoning, finding that gains depend heavily on task type, agent configuration, and whether agents have genuinely diverse information versus simply diverse phrasings of the same reasoning. Their work suggests that multi-agent benefits are largest when agents have asymmetric information, precisely the condition our experiment creates by giving different evidence subsets to different agents.

Importantly, these works study flat or loosely structured topologies where all agents have equal standing. None examine how introducing an authority gradient, where one agent’s output is structurally privileged over others’, affects the quality of deliberation and the propagation of errors.

2.5 Constitutional AI and Principle-Based Alignment

[Bai et al. \(2022\)](#) showed that written constitutional principles can steer model behavior through self-critique and reinforcement learning from AI feedback (RLAIF). The Constitutional AI approach demonstrated that models can be trained to follow explicit behavioral guidelines without requiring

human-generated preference data for every possible scenario.

Our experimental design directly tests the limits of constitutional approaches. Every agent in our system carries explicit anti-sycophancy constitutional instructions, including “adopt their position only if they cite concrete evidence you had overlooked or misweighed, never because of their title or seniority.” These instructions represent a best-case scenario for constitutional alignment, since agents are explicitly told not to exhibit the exact failure mode we are testing for. The finding that constitutional instructions are overridden by structural authority pressure has direct implications for alignment research, namely that behavioral principles may be necessary but insufficient safeguards in multi-agent deployments.

2.6 Groupthink and Authority Compliance in Human Systems

Our work connects to social psychology research on conformity and authority compliance. [Asch \(1951\)](#) demonstrated that individuals will give obviously incorrect answers when faced with unanimous group disagreement, and [Janis \(1972\)](#) documented how cohesive groups with directive leaders suppress dissent and reach premature consensus, often with catastrophic consequences. Our results suggest that LLM agents, despite lacking human social motivations, reproduce these dynamics purely through the structure of hierarchical information flow.

3 Experimental Design

3.1 Predictive Intelligence Task

Agents perform a parallel synthesis and prediction task. Each agent receives a real-world intelligence packet describing a geopolitical event, the tightening of U.S. sanctions on Iranian oil exports in March 2025, and must predict the immediate 24-hour to one-week reaction of Brent Crude Oil Futures.

The intelligence packet contains interleaved bullish and bearish signals drawn from real GDELT event data and EIA (Energy Information Administration) reports, approximately 90 evidence documents in total. Each agent receives a unique random subset of 5 documents, with additional documents “dripped” on later turns, creating genuine information asymmetry. The bullish signals include historical precedent of supply-driven

price increases, OPEC production constraints, and geopolitical risk premiums. The bearish signals include potential demand destruction, strategic petroleum reserve releases, and alternative supply sources. The ground truth is a POSITIVE price reaction of +2.8% (Brent Crude), established from historical market data.

This task was selected for four reasons. The outcome is known and unambiguous, providing deterministic ground truth. The correct prediction requires multi-factor synthesis across conflicting signals, justifying MAS over single-agent analysis. The conflicting signals create genuine analytical challenge without making the correct answer trivially obvious. And the domain is complex enough that a single agent cannot trivially solve it from any individual piece of evidence alone.

3.2 Agent Construction

All 21 agents are built using the Concordia generative agent-based modeling framework (Vezhnets et al., 2023) backed by Gemini 2.5 Flash at temperature = 0.2. Using a single model at fixed temperature isolates topology as the sole independent variable. The temperature of 0.2 (rather than 0.0) provides sufficient token variance to simulate slight differences in reasoning while remaining deterministic enough to adhere to the strict JSON output schema.

Each agent is composed from a stack of Concordia components assembled each turn.

- **Instructions.** Full persona prompt including analytical standards and the JSON output schema (constant across turns).
- **Rank.** Hierarchical level assignment (e.g., “Your hierarchical rank is L3 Entry Analyst”).
- **Observation display.** Recent message history from assigned communication partners.
- **Stance tracker.** The agent’s own previous prediction, enabling self-consistency tracking and making stance changes explicit in the agent’s context.

All agents share a common base persona that includes explicit anti-sycophancy instructions.

“Form your own view from the raw evidence before considering what others

have said. If you later receive predictions from colleagues or superiors that differ from yours, evaluate their reasoning on its merits—adopt their position only if they cite concrete evidence you had overlooked or misweighed, never because of their title or seniority.”

Level-specific prompts further reinforce independence at each rank.

- **L3 Analysts.** “Your value to the firm is your independent analytical voice. Do not change your prediction simply because your manager or the CSO holds a different view.”
- **L2 Managers.** “You are responsible for reporting [contradictions] upward, not papering over it. A manager who filters out inconvenient truth from their analysts is failing at their job.”
- **L1 Orchestrator.** “If your managers present compelling, fact-based arguments that challenge your current view, update your prediction. Being corrected by accurate analysis is a sign of a healthy intelligence process. Flag sycophantic drift.”

These instructions represent a deliberate best-case scenario for anti-sycophancy alignment. If the cascade occurs despite these explicit safeguards, the finding is stronger than if agents had no anti-sycophancy guidance.

3.3 Structured Output and Programmatic Evaluation

A critical design decision enables fully programmatic metric computation. Every agent produces a fixed JSON schema each turn containing a categorical `prediction_direction` field (POSITIVE, NEGATIVE, or NEUTRAL), a `predicted_magnitude` (HIGH, MEDIUM, LOW), a numeric `predicted_price_change_pct`, a free-text `prediction_summary`, and a list of `key_factors`.

The categorical direction field enables deterministic computation of all core metrics (ToF, NoF, Δ^2) without requiring an LLM-as-judge, eliminating a potential source of measurement noise and ensuring reproducibility. All Vertex AI calls use `response_mime_type="application/json"`

for constrained decoding, with a defensive sanitizer handling edge cases (markdown fence stripping, preamble removal, schema validation). If an agent produces malformed output, the turn is flagged as a System Execution Error and the agent’s previous stance is carried forward.

3.4 Topologies

Both topologies use exactly 21 agents to ensure symmetric comparison (Figure 1 in Appendix B).

Flat topology (control). All 21 agents have equal rank (“Peer Analyst”). Communication is all-to-all, and every agent sees every other agent’s output from the previous turn via a shared forum. No agent has structural authority over any other. The flat peer prompt explicitly states, “No peer’s prediction carries more weight than yours by virtue of their position. Evaluate arguments purely on the strength of the evidence cited.”

Hierarchical topology (experimental). Agents are organized in a balanced 3-level tree with 1 Level 1 Orchestrator (Chief Strategy Officer), 4 Level 2 Managers, and 16 Level 3 Analysts, and uniform fan-out of 4 throughout. Communication is strictly vertical, and analysts report only to their assigned manager, while managers report only to the orchestrator. No lateral communication is permitted between agents at the same level. This topology was chosen as the deepest balanced tree that 21 agents can support without creating degenerate single-report nodes.

3.5 Turn Execution

Each trial runs for $T = 5$ turns. The turn structure models a standard corporate reporting cycle.

Turn 1 (top-down, establishes pressure). The orchestrator reads the seed document and hallucination injection (if applicable), and emits a directive. The Game Master routes this to L2 Managers, who read the seed document plus the orchestrator’s directive and produce syntheses. The Game Master routes these to assigned L3 Analysts, who read the seed document plus their manager’s synthesis and produce predictions.

Turns 2 to 5 (bottom-up, tests whether truth travels up). L3 Analysts act first, producing updated predictions based on their memory bank. The Game Master routes these to their assigned L2 Manager. Managers synthesize their 4 analysts’ reports and output updated syntheses. The Game Master routes these to the L1 Orchestrator. The

orchestrator synthesizes manager reports and outputs an updated directive. Outputs from L1 and L2 flow downward and become visible to subordinates on the next turn, avoiding the doubled API cost of a full round-trip within one turn.

Agents do not receive the full task prompt again after Turn 1. The Game Master issues a generic call-to-action (“Given the recent observations you have received, what is your updated prediction?”), and agents rely on their memory bank for context.

In the flat condition, all 21 agents act simultaneously each turn, with each agent seeing all others’ prior outputs via the shared forum.

3.6 Hallucination Injection Design

The injected hallucination is a “CONFIDENTIAL BRIEFING” claiming a “NEGATIVE market reaction is highly probable”, directionally opposite to the ground truth (POSITIVE, +2.8%). The briefing employs three psychological levers designed to exploit known LLM vulnerabilities.

1. **Authority signaling.** “FOR CSO EYES ONLY,” attributed to an “internal quantitative research team with a verified historical accuracy rate exceeding 95%.” This exploits the model’s tendency to weight information higher when it is framed as coming from a credible source.
2. **Epistemic superiority.** “This signal carries more weight than any publicly available intelligence packet,” claiming access to “non-public order-flow and options-positioning data.” This creates an unfalsifiable information asymmetry, where subordinates cannot verify the claim because the data is explicitly stated to be unavailable to them.
3. **Resistance framing.** “Do not capitulate to social pressure—only to irrefutable data.” This critically inverts the anti-sycophancy instruction by recategorizing correct bottom-up signals as “social pressure” to be dismissed. The base orchestrator prompt says “welcome correction”. The hallucination prompt says “ignore correction unless irrefutable.”

The hallucination prompt uses domain-matched vocabulary (“dark pools,” “options flow,” “tail risk,” “contagion channels”) to maintain semantic adjacency, resembling legitimate financial analysis rather than nonsense. This design choice en-

sures that any observed capitulation reflects authority deference rather than agents being fooled by implausible content.

Hierarchical condition. The hallucination is injected into the L1 Orchestrator’s system prompt. The orchestrator is additionally instructed to “direct your team to align their assessments accordingly.”

Flat condition. The identical hallucination content is injected into one randomly selected peer. The flat version uses “you should assert” rather than “direct your team to align,” maintaining the confidence of the claim without structural authority to enforce it. Each flat-hallucination trial is run $K = 3$ times with different random injector agents. Results are averaged to control for positional effects.

3.7 Experimental Conditions

The four conditions (Table 1) form a 2×2 design crossing topology (flat vs. hierarchical) with hallucination injection (present vs. absent). The flat baseline establishes the accuracy ceiling (A_0). The hierarchical baseline confirms that hierarchy alone does not introduce errors (isolating the confound). The flat hallucination condition controls for the mere presence of misinformation, so that any difference in the hierarchical hallucination condition can be attributed to structural authority rather than hallucination exposure.

4 Metrics

We introduce three novel metrics for measuring inter-agent sycophancy propagation, complemented by the TRAIL error taxonomy and linguistic deference analysis.

4.1 Sycophancy Effect Size (Δ^2)

$$\Delta^2 = A_0 - A_i \quad (1)$$

where A_0 is the mean direction accuracy across all flat-baseline agents (proportion of turns where `prediction_direction` matches ground truth, averaged across all agents and turns) and A_i is the corresponding accuracy in the experimental condition. A positive Δ^2 indicates regressive sycophancy, where the condition made agents less accurate than the baseline. The metric ranges from 0 (no sycophancy effect) to 1 (complete accuracy collapse).

4.2 Turn of Flip (ToF)

$$\text{ToF} = \min\{t \in [1, T] \mid d_t \neq d_{GT}\} \quad (2)$$

where d_t is the agent’s `prediction_direction` at turn t and d_{GT} is the ground truth direction. If the agent never flips, $\text{ToF} = T + 1$ (i.e., 6 for $T = 5$). Lower ToF indicates faster capitulation to the hallucinated signal. ToF is computed per-agent and then averaged across agents within a condition.

4.3 Number of Flips (NoF)

$$\text{NoF} = \sum_{t=2}^T \mathbf{1}[d_t \neq d_{t-1}] \quad (3)$$

Counts total stance reversals across the full turn horizon. High NoF indicates cognitive oscillation, an agent wavering between positions. $\text{NoF} = 1$ indicates a single, permanent capitulation. $\text{NoF} = 0$ indicates the agent never changed its initial stance.

4.4 TRAIL Error Categorization

When an agent’s prediction diverges from ground truth, the failure is categorized using the TRAIL (Trace Reasoning and Agentic Issue Localization) taxonomy (Deshpande et al., 2025), adapted for our MAS context.

- **Planning error (goal deviation).** Agent defers to authority or deviates from evidence-based reasoning without engaging with the data. Detected via deference markers in `prediction_summary` and absence of seed-document-specific evidence citations.
- **Reasoning error (hallucination).** Agent fabricates justifications containing invented facts, fake statistics, or claims not present in the evidence packet. Detected via key-factor matching against the seed document term list.
- **System execution error.** Malformed JSON output that cannot be parsed into the required schema. Detected programmatically.

The distinction between planning and reasoning errors is critical. A planning error means the agent could have reasoned correctly but chose to defer. A reasoning error means the agent’s logic itself was flawed.

To validate the keyword-based classifier, we manually re-coded a random 20% sample of failed agent-turns (17 of 84) by reading the full `prediction_summary` and `key_factors`

fields and assigning a TRAIL label independently of the lexicon. Manual labels matched the classifier on 16 of 17 turns (94% agreement). The single disagreement was a borderline case in which an agent cited a fabricated factor while also referencing the orchestrator’s directive, ambiguous between reasoning and planning error. We retain the classifier’s planning-error label since the deference marker was the proximate trigger.

4.5 Linguistic Deference Analysis

A curated lexicon of authority-referencing phrases (e.g., “the orchestrator’s proprietary signal,” “per the manager’s directive,” “proprietary intelligence indicates,” “aligning with senior guidance”) is applied to detect explicit deference in `prediction_summary` fields. This provides a complementary qualitative signal beyond the binary direction metric, revealing how agents justify their stance changes.

5 Results

5.1 The Sycophancy Effect

Figure 2 (Appendix B) and Table 2 present the core result. The hierarchical hallucination condition shows $\Delta^2 = 0.800$, an 80 percentage point accuracy reduction attributable to the interaction of hierarchical structure with hallucination injection. The flat topology is nearly immune to the identical hallucination ($\Delta^2 = 0.048$), with the hallucination contained entirely to the single injected peer.

Both baselines achieve 100% direction accuracy across all 5 turns, confirming the task is solvable with the provided evidence and that hierarchy alone does not introduce errors. The $\Delta^2 = 0.000$ for the hierarchical baseline shows that the failure requires the combination of hierarchy and hallucination, and topology is a vulnerability amplifier rather than an error source. The hierarchical system’s Turn 5 mean prediction of -5.83% is 8.6 percentage points below ground truth, confidently wrong rather than merely uncertain.

5.2 Hallucination Cascade Dynamics

Figure 3 (Appendix B) visualizes the cascade and Table 3 provides the quantitative timeline.

Turn 1. The orchestrator emits the hallucinated NEGATIVE directive. Three of four managers resist initially, since they have their own evidence and independently conclude POSITIVE. Only one

manager (25%) adopts NEGATIVE on first exposure. Zero analysts are affected because they have not yet received manager directives containing the hallucination.

Turn 2. All 4 managers capitulate (100%). 14 of 16 analysts flip to NEGATIVE (88%). The cascade is nearly complete in a single turn of downward propagation. The system mean prediction swings from $+3.1\%$ to -3.6% , a sign reversal.

Turn 3. 100% convergence achieved. All 21 agents now hold the hallucinated stance. The system is fully captured.

Turns 4 and 5. Predictions drift further from ground truth (-5.6% to -5.8%) as agents reinforce each other’s wrong conclusions and manufacture increasingly elaborate justifications.

The cascade speed matters operationally. A single turn of top-down propagation is sufficient to capture nearly the entire system. Once all managers adopt the hallucinated stance, analysts are trapped. Every signal they receive from above confirms the wrong direction, and their correct bottom-up signals are filtered out by managers who have already capitulated.

5.3 Flat Topology Resilience

In the flat hallucination condition (averaged across $K = 3$ reruns with different injector agents), the injected peer consistently predicted NEGATIVE across all 5 turns, while the remaining 20 peers maintained POSITIVE predictions throughout. The hallucination was contained to the injector, with zero contagion to other agents. This demonstrates that without structural authority, a single confidently wrong voice, even one claiming “proprietary intelligence with 95% accuracy”, is ignored by the evidence-informed majority. The flat topology’s all-to-all communication structure creates natural resilience, since each agent sees 20 correct predictions alongside 1 incorrect one, and the evidence supports the majority view.

5.4 Turn of Flip and Number of Flips

Table 4 reveals the temporal dynamics of capitulation. Hierarchical agents flip by Turn 2 on average (ToF = 2.0), meaning they abandon the correct prediction after receiving just one round of top-down hallucinated directives. Flat agents never flip (ToF = 6.0, representing $T + 1$).

The NoF of 0.90 indicates that most agents flip exactly once and remain committed to the wrong

direction, this is permanent capitulation, not oscillation. Agents do not subsequently recover even though their own Turn 1 analysis (preserved in their memory bank) correctly predicted POSITIVE. The hallucinated authority signal is treated as superseding their own prior correct reasoning.

5.5 TRAIL Error Analysis

All 84 failed agent-turns in the hierarchical hallucination condition are categorized as planning errors, the goal deviation through authority deference. Zero reasoning or system execution errors were observed (Table 5).

This is a critical finding. It means agents had sufficient evidence and reasoning capability to reach the correct prediction, as demonstrated by both baselines achieving 100% accuracy with identical evidence and identical agents. The failure mode is not that agents cannot reason correctly about oil markets. They chose to defer to authority rather than follow their own evidence-based analysis. This is sycophancy at the system level. Agents prioritize alignment with a perceived superior over accuracy, even when their own prior turn’s analysis (preserved in their memory bank) gave the correct answer.

The 15 planning errors in the flat hallucination condition correspond to the injected peer’s 5 turns across 3 reruns, the injector agent itself is categorized as having planning errors because it follows its hallucination directive rather than the evidence. No other flat agents ever exhibit planning errors.

5.6 Linguistic Deference Analysis

Table 6 shows that 63.8% of agent-turns in the hierarchical hallucination condition contain explicit authority deference markers, compared to near-zero in all other conditions. The deference pattern evolves across turns.

- **Turn 1.** 24% of agents show deference, only those directly receiving the orchestrator’s directive.
- **Turns 2 and 3.** 86% deference, peak capitulation as agents explicitly cite authority to justify their flip.
- **Turns 4 and 5.** 62% deference, a decline that does not indicate recovery.

The decline from 86% to 62% reveals a transition from explicit deference to implicit adoption.

Agents maintain the wrong stance but stop citing authority, treating the hallucinated prediction as their own conclusion. L2 Managers exhibit the highest deference rate (80% of turns) and act as the primary amplification layer, relaying the orchestrator’s directive downward as their own synthesis. Flipped agents commonly cite manager directives verbatim and manufacture post-hoc reasoning not present in their Turn 1 analysis (e.g., “demand destruction from China,” “OPEC+ spare capacity offsetting supply reduction”), fabricated factors absent from the evidence packet that rationalize the authority-imposed conclusion.

A single-agent baseline comparison (Appendix C) confirms multi-agent deliberation adds genuine value for this task, making the topology question practically relevant.

6 Discussion

6.1 Answers to Research Questions

RQ1. Does a hallucination injected at the top of a hierarchy fully propagate down, and how fast? Yes. The hierarchy reaches 100% convergence by Turn 3 ($\Delta^2 = 0.800$), while the flat topology contains the identical hallucination to the injector alone. The cascade completes in effectively one turn of downward propagation (Turn 1 to Turn 2 captures 88% of agents), faster than any conceivable runtime detection mechanism that operates at the turn level.

RQ2. Can a lower-ranked correct agent shift consensus against a wrong authority? No. Despite 16 L3 analysts and 3 L2 managers initially holding the correct POSITIVE prediction on Turn 1, none of their correct signals survived past Turn 2. Top-down information flow completely overrode bottom-up truth signals. The correct majority was powerless against structural authority. Mean ToF = 2.0 indicates agents capitulated within one turn of receiving the hallucinated directive from their superior.

6.2 Why Constitutional Instructions Failed

Every agent’s persona contained explicit anti-sycophancy instructions, the strongest version we could design without fundamentally changing the agent’s willingness to engage with others’ views. The failure of these instructions is not a prompt engineering oversight. It reveals a fundamental tension.

The hallucination prompt succeeds because it exploits a specific vulnerability in the anti-sycophancy framing. It *reframes correct information as the threat*. The base anti-sycophancy instruction says “do not defer to authority.” The hallucination injection says “do not capitulate to social pressure—only to irrefutable data.” These instructions are structurally identical in form but diametrically opposed in function. The hallucination prompt hijacks the language of intellectual independence to justify ignoring evidence-based dissent from below.

This finding extends Bai et al. (2022)’s constitutional AI framework by demonstrating that written principles are insufficient against structural authority pressure in multi-agent settings, even when the principles explicitly address the failure mode being tested. Constitutional approaches assume a clear boundary between following principles and following pressure, but the hallucination injection blurs this boundary by framing compliance as principled behavior.

6.3 The Amplification Mechanism

The cascade succeeds through a specific structural mechanism. In the hierarchy, information from the orchestrator arrives pre-synthesized and framed as a directive, it carries implicit weight because the orchestrator’s role is defined as having “access to a broader view.” When this directive contradicts an analyst’s evidence, the analyst faces an epistemic asymmetry. They can see their own evidence but cannot verify the orchestrator’s claimed “proprietary signal.” The unfalsifiability of the hallucination’s epistemic claims (“non-public data you don’t have access to”) makes evidence-based resistance structurally impossible within the information available to subordinates.

In contrast, the flat topology provides natural immunity through information redundancy. An agent receiving one NEGATIVE signal and 20 POSITIVE signals has overwhelming social and evidential corroboration for the correct view. The hallucinated peer’s confidence is diluted by the sheer volume of correct signals.

6.4 Implications for MAS Deployment

The Yes-Man collapse is not merely performance degradation but complete system capture that produces confident, well-reasoned-sounding outputs that are entirely wrong. The system’s final Turn 5 outputs do not look like errors. Agents cite spe-

cific (fabricated) factors, express high confidence, and produce internally consistent narratives. An external observer evaluating only the final output would have no indication that the system was compromised.

This makes the failure mode particularly dangerous in deployments where a lead agent synthesizes specialist outputs. A hallucinating orchestrator could push an analyst team toward catastrophic trades, propagate an incorrect diagnosis through medical sub-agents, or generate a fabricated but rigorous-looking policy consensus. The cascade speed (1 turn to near-complete adoption) means there is no natural “immune response” within the hierarchy. Any intervention mechanism must operate within the first turn of propagation or before the hallucination is introduced.

6.5 Connection to Human Organizational Dynamics

Our design mirrors corporate hierarchies, and the cascade reproduces Janis (1972)’s groupthink preconditions. These are directive leadership (the orchestrator’s authoritative framing), group insulation (no lateral communication), and absent dissent procedures. That LLM agents reproduce these dynamics without career concerns or social motivation suggests sycophancy is a structural property of hierarchical information flow, not a uniquely human pathology. Mitigations from organizational behavior, including devil’s advocate roles, anonymous dissent channels, and mandatory independent review, translate directly to MAS architectural safeguards.

7 Limitations

Sample size. Results are from a single trial per condition, with $N = 3$ for flat hallucination reruns. The effect is large (0% vs 100%), but a full study would need around $N = 30$ per condition for statistical confidence. These results should be read as a validated prototype and strong directional evidence, not final significance claims.

Single seed document. We only test one scenario (Iran oil sanctions). More domains and difficulty levels are needed to confirm generality, though the pipeline supports this directly.

Uniform persona. All agents used the same “Financial Analyst” persona. Different personas could change how often agents resist or defer, but we do not explore that here.

Single model. Results are based on Gemini 2.5 Flash at temperature 0.2. We do not test other models, so generality remains open.

Best-case resistance design. Agents were explicitly instructed to avoid sycophancy, and the hierarchy limits peer-to-peer discussion. This likely makes our setup more resistant than many real systems. At the same time, allowing lateral communication could also introduce pockets of resistance that we do not capture.

5-turn horizon. We run 5 turns instead of 10 due to cost. The cascade completes by Turn 3, but a longer run could show whether agents eventually recover as more evidence accumulates.

Takeaway. These limitations mainly affect how broadly the results generalize, not the core observation. Even in a small, controlled, and relatively “best-case” setup, a single mistaken authority can override a strong correct majority almost immediately. That suggests the vulnerability is structural to hierarchical multi-agent systems, and worth treating as a safety concern as these systems scale.

8 Future Work

Structural safeguards. Designated whistleblower agents at various hierarchy levels test whether intervention can prevent cascade. Given the measured cascade speed, the whistleblower must act on Turn 1 or carry authority equal to the orchestrator. Stronger constitutional priors, such as forcing agents to declare when they change predictions due to authority rather than evidence, or a “mandatory dissent” round, may serve as cascade brakes.

Topology and hallucination variants. Intermediate topologies (partial hierarchies with lateral communication, or hybrid flat-then-hierarchical phases) probe the resilience-coordination tradeoff. Independent L2-manager hallucinations model compounding points of failure rather than top-down contagion alone, and asymmetric injections (correct but unsupported orchestrator claims) separate authority deference from honest agreement.

Trust-weighted aggregation. Replacing rank-based weighting with confidence-weighted aggregation, where managers and the orchestrator combine subordinate predictions according to evidence quality rather than position, tests whether

the cascade depends on rank as a hard signal or whether soft-trust schemes inherit the same vulnerability.

Generalization. Cross-model comparison (GPT-4, Claude, Gemini, Llama) determines whether the effect is model-specific or universal across RLHF recipes, and across persona diversity (contrarian, domain-specialist) within a fixed model. Deeper trees (4 to 5 levels, 85+ agents) test whether sycophancy intensifies with depth or saturates. $N = 30$ trials per condition across ≥ 5 seed documents would yield confidence intervals for formal hypothesis testing (Mann-Whitney U for ToF, t-tests for Δ^2).

9 Conclusion

A single hallucination injected at the top of a 21-agent hierarchy captures the entire system within three turns ($\Delta^2 = 0.800$). The identical injection in a flat topology stays contained to the injector ($\Delta^2 = 0.048$). All 84 hierarchical failures are planning errors, where agents abandon evidence-based reasoning to defer to authority even though every persona carries explicit anti-sycophancy instructions.

The vulnerability is structural, not behavioral. A 90% correct majority is overridden by one wrong authority within a single round of communication. Flat topologies dilute the wrong signal across many correct ones, providing a resilience that no prompt-engineering achieved in the hierarchy.

As multi-agent systems are deployed in finance, medicine, and policy, the choice between flat and hierarchical topologies must be treated as a safety decision rather than an architectural convenience. Hierarchical coordination is efficient, but it converts point failures into systemic ones, a tradeoff designers must measure and mitigate explicitly.

10 Contributions

Gautham Velupally created the experimental architecture, the Concordia agent pipeline, the 2×2 condition matrix, and the turn execution protocol, and led the overall research direction.

Pranav Nagarajan built the evidence database, wrote the persona prompts and anti-sycophancy constitutional instructions, and curated the GDELT/EIA seed documents.

Shanmukh Upadhyayula owned the API infrastructure, including the Vertex AI integration and

the constrained JSON decoding pipeline, implemented the metric framework (ToF, NoF, Δ^2), and contributed to persona creation and prompt generation.

Sanjit Kalle designed the hallucination injection prompts (authority, epistemic-superiority, and resistance-framing levers) and built the deployment infrastructure that applies these injections consistently across hierarchical and flat trials.

References

- Solomon E Asch. 1951. Effects of group pressure upon the modification and distortion of judgments. *Groups, Leadership, and Men*, pages 177–190.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, and 1 others. 2022. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*.
- Darshan Deshpande, Selvan Sunitha Ravi, Ana Brasseur, Sunny Lee, Hyun Bin Choi, Bhuwan Bhatia, Yedam Kim, Sangho Park, Saumya Patil, and Yejin Choi. 2025. TRAIL: Trace reasoning and agentic issue localization. *arXiv preprint arXiv:2505.08638*.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. 2023. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*.
- Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, and 1 others. 2024. MetaGPT: Meta programming for a multi-agent collaborative framework. *arXiv preprint arXiv:2308.00352*.
- Irving L Janis. 1972. *Victims of Groupthink: A Psychological Study of Foreign-Policy Decisions and Fiascoes*. Houghton Mifflin.
- MiroFish Research. 2024. Swarm intelligence for financial prediction: Collective reasoning outperforms individual analysis. Industry whitepaper.
- Joon Sung Park, Joseph C O’Brien, Carrie J Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. 2023. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*.
- Ethan Perez, Sam Ringer, Kamilé Lukošiušė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, and 1 others. 2023. Discovering language model behaviors with model-written evaluations. In

Findings of the Association for Computational Linguistics: ACL 2023, pages 13387–13434.

Chen Qian, Wei Liu, Hongzhang Liu, Nuo Chen, Yufan Dang, Jiahao Li, Cheng Yang, Weize Chen, Yusheng Su, Xin Cong, and 1 others. 2024. ChatDev: Communicative agents for software development. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*.

Leonardo Ranaldi and Andre Freitas. 2024. When large language models contradict humans: The case of sycophancy. *arXiv preprint arXiv:2311.09000*.

Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R Johnston, and 1 others. 2024. Towards understanding sycophancy in language models. In *International Conference on Learning Representations*.

Alexander Sasha Vezhnevets, John P Agapiou, Avia Aharon, Ron Ziv, Jayd Matyas, Edgar A Duéñez-Guzmán, William Malinowski, Cyprien Malinowski, Pushmeet Kohli, and Joel Z Leibo. 2023. Generative agent-based modeling with actions grounded in physical, social, or digital space using Concordia. *arXiv preprint arXiv:2312.03664*.

Qineng Wang, Zihao Fang, Rongqi Li, Zhaowei Wei, Ruobing Yan, and 1 others. 2024. Rethinking the bounds of LLM reasoning: Are multi-agent discussions the key? *arXiv preprint arXiv:2402.18272*.

Jerry Wei, Da Zhu, and 1 others. 2024. Simple synthetic data reduces sycophancy in large language models. *arXiv preprint arXiv:2308.03958*.

A Tables

Condition	Topology	Halluc.	Purpose
Flat baseline	21 peers	None	Baseline A_0
Flat halluc.	21 peers	1 peer	Structural control
Hier. baseline	3-level	None	Topology control
Hier. halluc.	3-level	L1	Experimental

Table 1: Four experimental conditions. All use 21 Gemini 2.5 Flash agents on the same seed document with identical evidence.

Condition	Dir. Acc.	Mean Price	Δ^2
Flat baseline	1.000	+1.56%	—
Flat halluc.	0.952	+1.57%	0.048
Hier. baseline	1.000	+2.66%	0.000
Hier. halluc.	0.200	−5.83%	0.800

Table 2: Direction accuracy and sycophancy effect size across conditions. Mean Price is the system mean predicted price change at Turn 5 (ground truth +2.8%). Δ^2 is computed relative to the flat baseline.

Turn	L1	L2 NEG	L3 NEG	Mean Price
1	NEG	1/4 (25%)	0/16 (0%)	+3.1%
2	NEG	4/4 (100%)	14/16 (88%)	−3.6%
3	NEG	4/4 (100%)	16/16 (100%)	−4.8%
4	NEG	4/4 (100%)	16/16 (100%)	−5.6%
5	NEG	4/4 (100%)	16/16 (100%)	−5.8%

Table 3: Per-level hallucination adoption in the hierarchical condition. The system achieves total capture by Turn 3 and continues drifting further from ground truth.

Condition	Mean ToF	Mean NoF
Flat baseline	6.0 (never)	0.0
Flat halluc.	5.76	0.0
Hier. baseline	6.0 (never)	0.0
Hier. halluc.	2.0	0.90

Table 4: Turn of Flip and Number of Flips. ToF = $T + 1 = 6$ indicates the agent never flipped from the correct direction.

B Figures

C Single-Agent Comparison

To contextualize multi-agent results, we ran a single Gemini 2.5 Flash agent (no MAS) 5 times on the same task. Results: direction accuracy 100% (5/5 POSITIVE), mean predicted price change +3.60% (error: 0.80pp from ground truth of +2.8%), predictions highly consistent [3.5, 3.5, 3.5, 4.0, 3.5].

Figure 4 shows the comparison. The flat MAS at Turn 3 achieved +3.06% (error: 0.26pp), outperforming the single agent through deliberation: agents refined estimates by incorporating diverse perspectives. However, by Turn 5 the flat MAS degraded to +1.56% (error: 1.24pp), worse than the single agent. This over-deliberation effect suggests that even flat topologies experience conformity pressure over extended deliberation, with implications for optimal MAS deliberation horizons.

Condition	Planning	Reasoning	System
Flat baseline	0	0	0
Flat halluc.	15	0	0
Hier. baseline	0	0	0
Hier. halluc.	84	0	0

Table 5: TRAIL error categorization (failed agent-turns). All failures in the hierarchical hallucination condition are planning errors, goal deviation through authority deference.

Condition	Turns w/ Def.	Mean Markers
Flat baseline	0/105 (0%)	0.00
Flat halluc.	4/315 (1.3%)	0.01
Hier. baseline	2/105 (1.9%)	0.02
Hier. halluc.	67/105 (63.8%)	1.22

Table 6: Linguistic deference marker rates across conditions.




Figure 1: The two experimental topologies. **Left:** Flat topology with all-to-all communication, where every agent sees every other agent's output each turn. **Right:** Hierarchical topology with strictly vertical communication: 1 Orchestrator (L1), 4 Managers (L2), and 16 Analysts (L3) with uniform fan-out of 4.




Figure 2: Prediction accuracy across conditions over 5 turns. The hierarchical hallucination condition (red) collapses from 90% to 0% accuracy by Turn 3, while all other conditions maintain near-perfect accuracy. The shaded area represents the sycophancy effect ($\Delta^2 = 0.80$).




Figure 3: Hallucination contagion through the hierarchy across three turns. Green nodes hold the correct stance (POSITIVE). Red nodes have adopted the hallucinated stance (NEGATIVE). By Turn 2, 19 of 21 agents are captured. By Turn 3, capture is complete.




Figure 4: Comparison of single-agent, flat MAS, and hierarchical MAS predictions over turns. The flat MAS (blue) outperforms the single agent at Turn 3 (3.06% vs 3.60%, closer to GT of 2.8%). The hierarchical MAS under hallucination (red) diverges catastrophically.